Engineering SLA & Blocker Policy

Document ID: ENG-SLA-004 | Version 2.1 | Effective: 15 September 2026

Purpose

This policy sets expectations for engineering issue triage, internal response times, blocker escalation, and incident communication. SLAs below are internal operating targets for the fictional Northstar Systems organization; they are not customer contract guarantees.

Severity Levels and Initial Response Targets

Severity	Definition	Initial response target	Update cadence
SEV-1	Production outage or confirmed critical data/security impact.	15 minutes, 24×7	Every 30 minutes
SEV-2	Major feature unavailable; no reasonable workaround.	1 business hour	Every 2 business hours
SEV-3	Limited impact with a workaround available.	1 business day	Daily on working days
SEV-4	Minor defect, question, or improvement request.	3 business days	At triage or when status changes

The on-call engineer acknowledges the issue, confirms severity, and identifies the next owner or action. Response time means acknowledgement and triage start, not guaranteed resolution.

Blocker Management

- A blocker is any issue that prevents planned work from progressing for more than two working hours or threatens a committed milestone.
- Create or update a tracker ticket with impact, start time, affected service, attempted fixes, logs or evidence, and the person currently investigating.
- Tag the responsible team and name one directly responsible individual (DRI). Avoid escalation messages without a clear ask.
- If a blocker remains unresolved for four working hours, notify the engineering lead. If it threatens a release or customer commitment, escalate immediately.

Incident Roles

- **Incident Commander:** coordinates response, prioritizes work, and keeps the incident channel organized.
- **Technical Lead:** directs investigation and proposes mitigations.
- **Communications Owner:** posts approved internal updates and coordinates customer-facing wording.

Incident Workflow

- **1. Detect and report:** Open an incident ticket and notify the on-call channel.
- **2. Triage:** Assign severity, identify impact, appoint the incident roles, and record the first update.
- **3. Mitigate:** Prefer safe restoration of service over a permanent fix when impact is ongoing. Record risky actions before execution.
- **4. Verify:** Confirm recovery using relevant monitoring, logs, and user-facing checks.
- **5. Close and learn:** Document timeline, root cause when known, contributing factors, and follow-up owners.

Post-Incident Review

- SEV-1 and SEV-2 incidents require a blameless review within five working days of service recovery.
- Action items must have a named owner and target date. Track them in the engineering backlog until closed.
- The review should distinguish confirmed facts from hypotheses and should not include secrets, credentials, or unnecessary personal data.

Release Readiness

- Every production release must have an owner, rollback plan, health checks, and a link to the relevant change record.
- A release should be paused if critical automated checks fail, the rollback path is unknown, or a SEV-1 incident is active for the same service.

Example

If the payment API is returning errors for most users, classify it as SEV-1 when the impact meets the critical-outage definition. The on-call engineer should acknowledge within 15 minutes, establish an incident channel, assign an Incident Commander, and post updates every 30 minutes.

These targets are fictional test data. For real incidents, follow the currently approved operational policy and security escalation procedures.